

Atomic and Fair Data Exchange via Blockchain

최원재

2025. 01. 04



Paper Info

Title

- Atomic and Fair Data Exchange via Blockchain
- [Link](#)

Conference


- CCS '24: Proceedings of the 2024 on ACM SIGSAC Conference on Computer and Communications Security

Authors

- Ertem Nusret Tas
 - Ph.D Candidate at Stanford
 - Also author of [Bitcoin-enhanced proof-of-stake security: Possibilities and impossibilities](#)
- A16z

Problem Statement

The Fair Data Exchange (FDE) Problem



Server Fairness

- An adversarial client cannot learn anything about the data without paying the server.

Client Fairness

- An adversarial server cannot receive any payment if the client does not obtain the data.

Correctness

- If the client and the server are honest, the client obtains the data, and the server obtains the payment.

Existing Solutions

Subscription-based

- Client pays the server in advance
- Trust the reputation of the server to deliver the data.
- No Client Fairness

Free Data

- Server provides data for free
- Lack of Incentive for the server.
- No data integrity
- No Server Fairness

Blockchain

- Too Expensive (1MB = \$3000)

Proof of Storage / Retreivability

- Only guarantees and incentivize the storage of data
- No guarantee and incentive for the real delivery of the data

FDE without Trusted Third Party is Impossible

➤ ['On the Impossibility of Fair Exchange without a Trusted Third Party', Pagnia and Gartner, 1999](#)

Blockchain as a Trusted Third Party!

Usecases

Private Data Exchange

- Private Key

Data Marketplace

- Movies
- AI Models

Proto-Danksharding (EIP-4844)

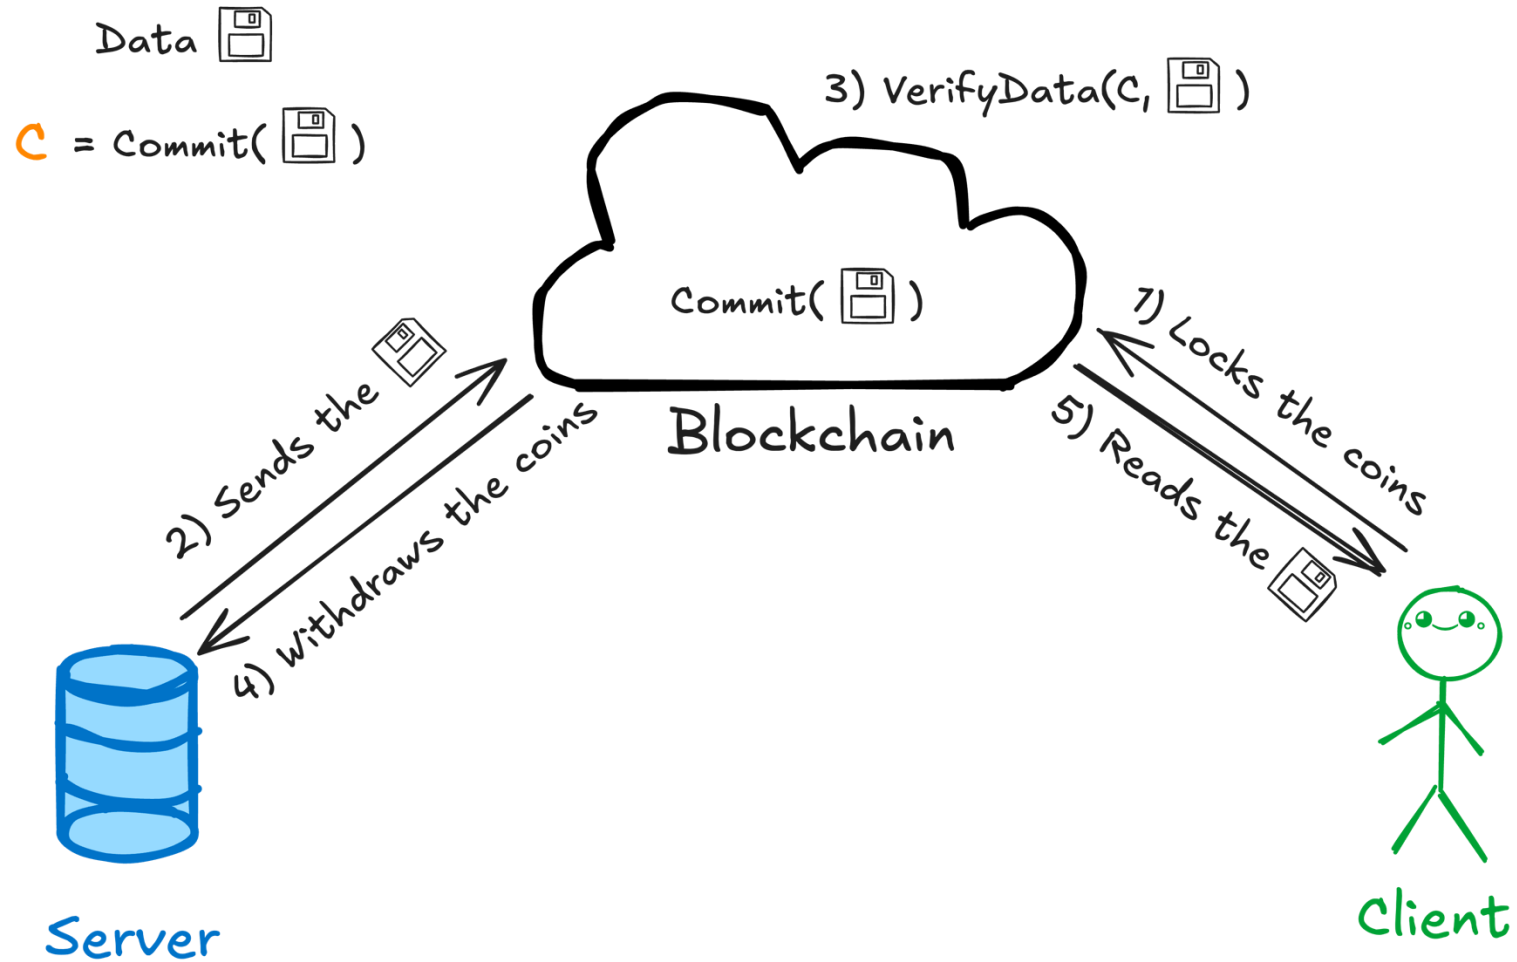
- Blob data is not stored on the Ethereum.
- Only the KZG Commitment is stored.
- Servers(Nodes) voluntarily store the real blob data and provides the the clients.
- FDE can incentivize it

Blockchain Archive Data

- Archive data is also not incentivized but needs big amount of data storage

FDE Protocol

Naïve Solution



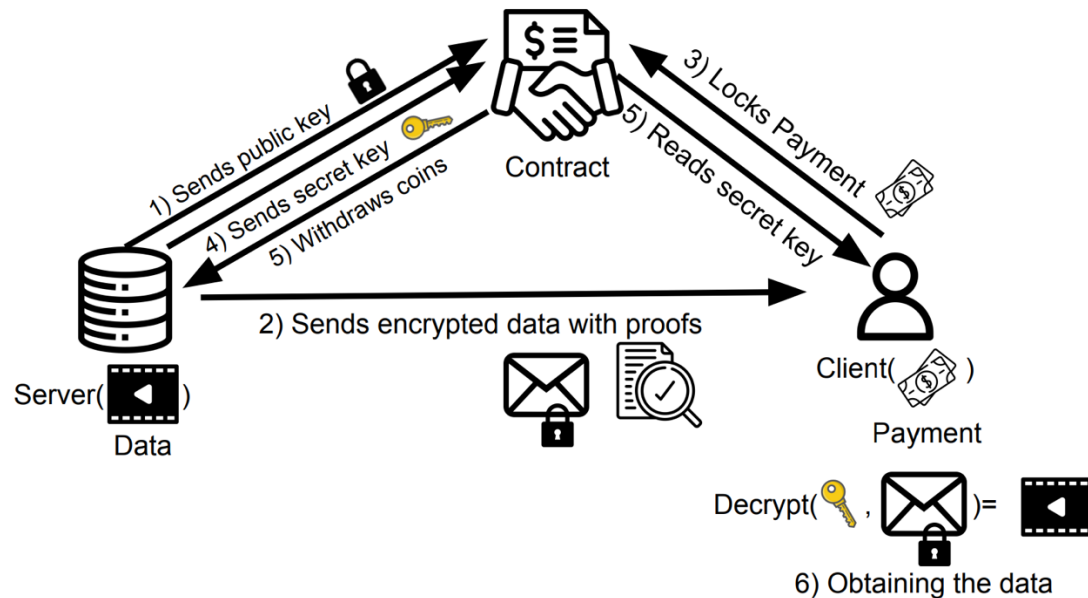
Problems

All data is posted to the blockchain

Two Problems

- Not only the buyer gets the data but everyone can get the data.
- Uploading all the data to the blockchain is impossible(or expensive) due to the limited capacity.

FDE Protocol



Solution

- Server encrypts the data
- Transfers the encrypted data off-chain
- Sells the decryption key on-chain

Must Prove

- The key decrypts the ciphertext to the data under C

Verifiable Encryption under Committed Key (VECK)

Common

➤ $Gen(crs) \rightarrow pp$

Server

➤ $Enc(C, w) \rightarrow (vk, sk, ct, \pi)$

- Proof π proves ct is the encryption of the data committed by C under sk committed by vk .

Blockchain

➤ $Verify_{key}(vk, sk) \rightarrow 1 \text{ or } 0$

Client

➤ $Verify_{ct}(C, vk, ct, \pi) \rightarrow 1 \text{ or } 0$

➤ $Decrypt(sk, ct) \rightarrow w \text{ or } \perp$

Properties of VECK

Correctness

- Verifications for honestly generated encryption succeed.
- Correctness of FDE

Soundness

- No PPT adversary can generate vk, sk, ct, π such that verification succeeds, yet decryption does not output w .
- Client Fairness of FDE

Computational Zero-Knowledge

- The ciphertext and the proof leak no additional information about the witness.
- Server Fairness of FDE

KZG Commitment

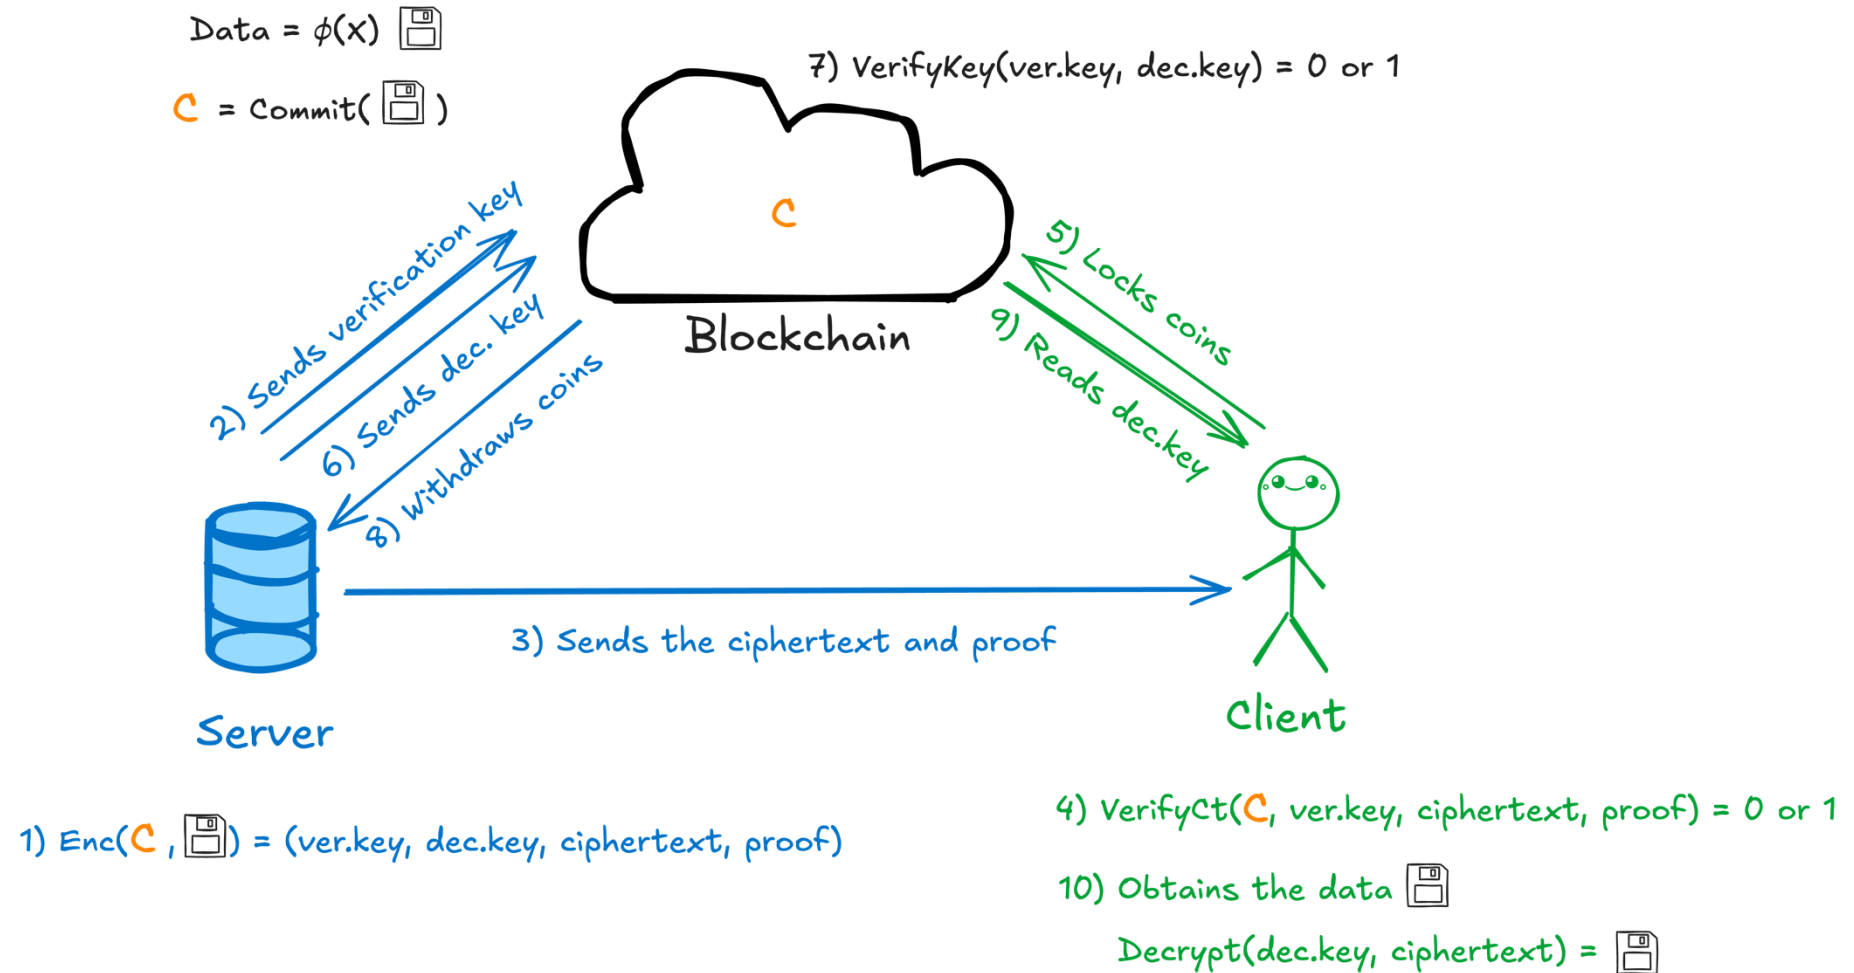
Proto-Danksharding (EIP-4844)

- For the utilization for proto-danksharding, the paper assumes the KZG Commitment scheme.

KZG Commitment

- $w = \phi$
- $data = (\phi(0), \dots, \phi(l))$
- $C = CommitKZG(\phi)$

FDE Protocol for KZG Commitments



Two Encryption Algorithms

ElGamal Encryption

Paillier Encryption

ElGamal Encryption

Key Generation

- p : Big Prime Number
- g : Primitive Root
- x : Random Sampled Secret Key
- y : Public Key
 $y = g^x \bmod p$

- Asymmetric Key Encryption Algorithm
- Based on the Diffie-Hellman Key Exchange
- Security depends on Decisional Diffie Hellman Problem
- Discrete Logarithm Problem

Encryption

- Ciphertext (c_1, c_2)
 - $c_1 = g^k \bmod p$
 - $c_2 = m \cdot y^k \bmod p$

Decryption

- $s = c_1^x \bmod p$
- $m = c_2 \cdot s^{-1} \bmod p$

ElGamal-based VECK (SNARK version)

Prover (Server)

$Enc(C, \phi(X)) \rightarrow (sk, vk, ct, \pi)$

- Sample $s \leftarrow_R \mathbb{F}_P$
- Set $sk := s$ (ElGamal decryption key)
- $vk := h^s$ (corresponding verification key)
- $ct := \{h_i^s g_i^{\phi(i)}\}_{i=0}^l$ (ElGamal ciphertexts)
- Generate $\pi := SNARK.Prove(crs, instance = (C, vk, ct), witness = (s, \phi(X)))$

Verifier (Client)

$Verify_{ct}(C, vk, ct, \pi) \rightarrow 0 \text{ or } 1$

- Output 1 if π verifies against C, vk, ct

$Decrypt(sk, ct) \rightarrow w \text{ or } \perp$

- Decryption algorithm for ElGamal

Blockchain Contract

$Verify_{key}(vk, sk) \rightarrow 0 \text{ or } 1$

- Check if $h^{sk} == vk$

zkSNARK is Inefficient!

Efficient VECK protocol

- In this paper, the protocol utilizes the shared structure of the ElGamal & KZG

Splitting the Message

- Can support large message by splitting the message into k chunks with range proofs.

Protocol Breakdown

VECK Protocol with ElGamal Encryption and KZG Commitment

GEN(crs) \rightarrow pp: On input crs = $(\mathcal{G}, \{g_1^{\tau^i}\}_{i=1}^n, \{g_2^{\tau^i}\}_{i=1}^n)$, generate random group elements with unknown discrete logarithms $h_i \leftarrow_R \mathbb{G}_1, h \leftarrow_R \mathbb{G}_1$ for $i \in [n]$. In practice, these can be generated using a Hash-to-Curve [9] function $H_{\text{curv}}: \mathbb{F}_p \rightarrow \mathbb{G}_1, h_i := H_{\text{curv}}(i)$. Output pp = $(\text{crs}, \{h_i\}_{i=0}^n, h)$.

ENC($F_{[\ell]}^{\text{full-eval}}, C_\phi, \phi(X)$) \rightarrow (vk, sk, ct, π) :

- (1) Sample $s \leftarrow_R \mathbb{F}_p$, set sk := s, vk = h^s .
- (2) Compute the encryptions $\{\text{ct}_i := h_i^s \cdot g_1^{\phi(i)}\}_{i=0}^\ell$.
- (3) Generate a Fiat-Shamir pseudo-random challenge $\alpha := H(C_\phi, \text{ct}_0, \dots, \text{ct}_\ell, [\ell])^a$.
- (4) Compute $C_\alpha := g_1^{\phi(\alpha) + s(\tau - \alpha)}$.
- (5) Compute the opening proof for α : $\pi_\alpha := \text{OPEN}(\text{crs}, \alpha, \phi(X) - s(X - \alpha))$ (note that $\pi_\alpha = g_1^{(\phi(\tau) - \phi(\alpha)) / (\tau - \alpha) - s}$).
- (6) Compute π_{LIN} to prove that C_α is indeed of the form g_1 raised to a known power, multiplied by $g_1^{\tau - \alpha}$ raised to the secret s of the verification key, i.e., we compute the proof for the relation \mathcal{R}_{LIN} as described in Appendix A.2, Equation (A.4):
$$\mathcal{R}_{\text{LIN}} = \left\{ \begin{array}{ccc|c} g_{1,1} = g_1 & g_{1,2} = g_1^{\tau - \alpha} & u_1 = C_\alpha & x_1 = \phi(\alpha) \\ g_{2,1} = 1 & g_{2,2} = h & u_2 = \text{vk} & x_2 = s \end{array} \mid g_{1,1}^{x_1} g_{1,2}^{x_2} = u_1 \wedge g_{2,1}^{x_1} g_{2,2}^{x_2} = u_2 \right\}.$$
- (7) Compute $Q := \prod_{i=0}^\ell h_i^{L_{i,\ell}(\alpha)} \cdot g_1^{-(\tau - \alpha)}$.
- (8) Compute discrete logarithm equality π_{DLeq} with respect to (Q, Q^s, h, vk) for the witness s, i.e., we compute the proof for the relation $\mathcal{R}_{\text{DLeq}}$ as described in Appendix A.2, Equation (A.3):
$$\mathcal{R}_{\text{DLeq}} := \{(g = Q, h, a = Q^s, b = \text{vk}); x = s \mid a = g^x \wedge b = h^x\}.$$
- (9) Output (sk, vk, ct = $[\text{ct}_0, \dots, \text{ct}_\ell]$, $\pi = (C_\alpha, \pi_\alpha, \pi_{\text{LIN}}, \pi_{\text{DLeq}})$).

VER_{ct}($F_{[\ell]}^{\text{full-eval}}, C_\phi, \text{vk}, \text{ct}, \pi$) \rightarrow 0/1 :

- (1) Parse ct = $(\text{ct}_0, \dots, \text{ct}_\ell)$ and the server's proofs $\pi = (C_\alpha, \pi_\alpha, \pi_{\text{LIN}}, \pi_{\text{DLeq}})$.
- (2) Verify π_{LIN} for C_α against $g_1, g_1^{\tau - \alpha}, h, C_\alpha, \text{vk}$.
- (3) Compute the Fiat-Shamir challenge: $\alpha := H(C_\phi, \text{ct}_0, \dots, \text{ct}_\ell)$.
- (4) Verify the opening proof π_α : $e(C_\phi / C_\alpha, g_2) = e(\pi_\alpha, g_2^{\tau - \alpha})$.
- (5) Compute Q as above for encryption.
- (6) Compute ct := $\prod_{i=0}^\ell \text{ct}_i^{L_{i,\ell}(\alpha)} \in \mathbb{G}_1$, and $Q^* := \text{ct} / C_\alpha$.
- (7) Verify π_{DLeq} for the relation (Q, Q^*, h, vk) .
- (8) If any of the verification checks fail, output 0. Otherwise, output 1.

VER_{key}(vk, sk) \rightarrow 0/1 : For sk = s $\in \mathbb{F}_p$, return 1 if and only if vk = h^s .

DEC($F_{[\ell]}^{\text{full-eval}}, \text{sk}, \text{ct}$) $\rightarrow \{\phi(i)\}_{i=0}^\ell$: For sk = s $\in \mathbb{F}_p$, for each $i \in [\ell]$, recover the value $x = \phi(i)$ by brute-forcing $x \in [\mathcal{B}]$ in $g_1^x = \text{ct}_i / h_i^s$. Pollard's rho algorithm for logarithms [59], Shank's baby-step-giant-step algorithm [65] or brute-force with pre-computed tables [22] could be helpful to find the right trade-off between decryption memory and time.

^aGeneralizing to more general subsets S other than $[\ell]$, the description of the subset S would be part of the input of $H(\cdot)$ generating the Fiat-Shamir challenge.

VECK Protocol with ElGamal Encryption and KZG Commitment

1. Setup

➤ Establish Reference

- Generates a common reference string (CRS) and public parameters.
- Aligns all participants on the same group-based environment for commitments, encryption and proof verification.

2. Encryption

➤ Key Generation

- Server picks a random exponent (secret key) and derives a corresponding verification key.

➤ Per-Point Ciphertexts

- Evaluates the polynomial at each point and encrypts these values via group exponentiation.

➤ Proofs of Correctness

- Produces short zero-knowledge proofs asserting that the ciphertexts match the committed polynomial under the same secret exponent.

High Level Explanation of the Proof

1. KZG Commitment

- The server has a **committed polynomial** $\phi(X)$ via a KZG commitment, $C\phi$.
- This commitment fixes the polynomial but does not reveal its coefficients.

2. Ciphertexts and a Secret Exponent

- For each point i , the server encrypts the value $\phi(i)$ using a secret exponent s .
- Each ciphertext ct_i encodes both $\phi(i)$ and the exponent s .

3. Random Challenge α

- A **Fiat–Shamir**-style hash function takes the public data (commitment, ciphertexts) and yields a random challenge α .
- This binds the proof to the specific commitment and ciphertext set, preventing replay or substitution attacks.

4. Blinded Evaluation at α

- The server “opens” (proves knowledge of) the polynomial at α , **but** adds a twist: it mixes in the same exponent s .
- Concretely, the server shows a short proof that
 - $\phi(\alpha)$ is consistent with the KZG commitment,
 - s is used in a linear combination with $\phi(\alpha)$, ensuring the exponent matches that used in the ciphertexts.

High Level Explanation of the Proof (Continued)

5. Aggregation via Lagrange Polynomials

- The ciphertexts $\{ct_i\}$ are cleverly combined (using Lagrange coefficients) to simulate an “encryption” of $\phi(\alpha)$.
- If that aggregated encryption matches the server’s **blinded** version of $\phi(\alpha)$ from the previous step, it confirms each ct_i was formed correctly.

6. Discrete-Log Equality (DLeq) Check

- Finally, there is a **discrete log equality proof** to confirm the server used **the same secret exponent** s across:
 - The KZG-related portion, and
 - The encryption process.
- If this DLeq check holds, it proves the exponent in the ciphertexts truly matches the exponent that “opens” the polynomial.

By combining these steps, the verifier sees that:

- The polynomial $\phi(X)$ in the KZG commitment is the **exact** polynomial the server used for all evaluations.
- Each ciphertext ct_i **truly** encrypts $\phi(i)$.
- A **single** secret exponent s ties everything together, so there is no mismatch or cheating between commitment and ciphertexts.

VECK Protocol with ElGamal Encryption and KZG Commitment

3. Verification of Ciphertexts

➤ Challenge Binding

- The verifier recomputes or checks a challenge (often via a hash) that links the polynomial commitment, the ciphertexts, and the server's claimed verification key.

➤ Check Consistency

- Verifies that each ciphertext matches the supposed polynomial evaluation and that all proofs pass. This involves checking if the same secret key was consistently used and if the commitments align with the encrypted values.

4. Verification of Key

➤ Server Reveals Secret Exponent

- At some stage (often to claim payment), the server discloses its secret exponent to a smart contract or directly to the client.

➤ Key-Value Matching

- A simple group-based check confirms that the revealed exponent truly corresponds to the previously published verification key. This step prevents a dishonest server from providing a key that does not actually decrypt the ciphertexts.

VECK Protocol with ElGamal Encryption and KZG Commitment

5. Decryption

➤ Recover Polynomial Evaluations

- The client (or any holder of the secret exponent) undoes the exponentiation in the ciphertexts to retrieve the original polynomial values.


➤ Discrete Log Lookup

- Depending on how the encryption was structured (e.g., exponential ElGamal), the client may perform a discrete-log search (e.g., baby-step–giant-step) to recover the numeric evaluation from the group element.

➤ Fair Exchange Outcome

- The client, having paid or locked tokens, now obtains the genuine data—fulfilling the exchange contract on-chain or off-chain with verifiable correctness and fairness.

Multiclient Model



Cannot reuse sk, vk, ct across clients

- If clients share the data, they can know without the payment
- Must rerandomize sk across clients!

Multi-client VECK Protocol

- Prover saves work by moving parts of the proof generation to a preprocessing step.
- No need to generate new ciphertexts and range proofs per client!

Evaluations


Performance Evaluation

On a Consumer PC Setting

- AMD Ryzen 5 3600 CPU and 8GB RAM

Prover time for 4096 (BLS12-381) field elements (128KB)

- ElGamal Encryption ($k=8$)
 - Range Proofs + Ciphertexts: 89s
 - Proving the consistency of ciphertexts: < 40ms
 - Ciphertexts & Proofs total size: 1.56MB
 - Proof size constant to the data
- Paillier Encryption
 - 5.09s
 - Ciphertexts & Proofs total size: 6.55MB
 - Proof size linear to the data



On-chain Cost

Transaction	Gas cost		USD cost	
	ElGamal	Paillier	ElG.	Pail.
serverSendsPubKey	158, 449	176, 296	5.11 \$	5.68 \$
clientLocksPayment	30, 521	30, 521	0.98 \$	0.98 \$
serverSendsSecKey	73, 692	82, 475	2.37 \$	2.65 \$
withdrawPayment	43, 836	43, 836	1.41 \$	1.41 \$

Table 2: EVM gas costs of the smart contract components for $\lambda = 128$ bits of security. Note these costs are constant in the size of the exchanged data, (cf. Appendix C.2), in the Exponential Elgamal (cf. Figure 2) and Paillier encryption-based FDE schemes (cf. Figure 5). For the USD costs, we used the transaction fee 14 GWei and 2, 302.35 USD/ETH exchange rate as of 2024 February 3rd.

Miscellaneous

Can be Implemented on Bitcoin

- No need for turing complete contract

Github

- [Link](#)

Thank You